

Hadoop Developer Training – Hive Lab Book

Table of Contents

Lab 1: Understanding Hive Tables	3
Lab 2: Analyzing Big Data with Hive	10
Lab 3: Understanding MapReduce in Hive.....	17
Lab 4: Joining Datasets in Hive.....	21
Lab 5: Computing ngrams of Emails in Avro Format....	25
Lab 7: Using HCatalog with Pig.....	32
Lab 8 : Advanced Hive Programming	37
Lab 9 : Streaming Data with Hive and Python.....	48

Location of Files:	/<userpath>/labs/
--------------------	-------------------

Note: Please use the below paths to avoid error while practicing

1. <userpath> should be your working directory(user home directory)
2. for Cloudera HDFS path is : user/training/
3. for Hartonworks HDFS path is : user/root/

Lab 1: Understanding Hive Tables

Perform the following

Step 1: Review the Data

1.1. Use the **ls** command to view the contents of the **wh_visits** folder. You should see 6 **part-m** files:

```
# hadoop fs -ls /apps/hive/warehouse/wh_visits/
```

1.2. Recall that the Pig projection to create these files had the following schema:

```
project_potus = FOREACH potus GENERATE
$0 AS lname:chararray,
$1 AS fname:chararray,
$6 AS time_of_arrival:chararray,
$11 AS appt_scheduled_time:chararray,
$21 AS location:chararray,
$25 AS comment:chararray ;
```

In this lab, you will define Hive table that matches these records and contains the exported data from your Pig script.

Step 2: Define a Hive Script

2.1. In the **Lab7.1** folder, there is a text file named **wh_visits.hive**. View its contents. Notice it defines a Hive table named **wh_visits** with the following schema that matches the data in your **project_potus** folder:

```
# more wh_visits.hive
create table wh_visits (
lname string, fname string, time_of_arrival string, appt_scheduled_time string,
meeting_location string, info_comment string)
ROW FORMAT DELIMITED
FIELDS TERMINATED BY '\t' ;
```

NOTE: You cannot use **comment** or **location** as column names because those are reserved Hive keywords, so we changed them slightly.

2.2. Run the script with the following command:

```
# hive -f wh_visits.hive
```

2.3. If successful, you should see “OK” in the output along with the time it took to run the query.

Step 3: Verify the Table Creation

3.1. Start the Hive Shell:

```
# hive hive>
```

3.2. From the **hive>** prompt, enter the “**show tables**” command:

```
hive> show tables;
```

You should see **wh_visits** in the list of tables.

OK		
lname	string	None
fname	string	None
time_of_arrival	string	None
appt_scheduled_time	string	None
meeting_location	string	None
info_comment	string	None

3.3. Use the **describe** command to view the details of **wh_visits**:

```
hive> describe wh_visits;
```

3.4. Try running a query (even though the table is empty):

```
select * from wh_visits limit 20;
```

You should see 20 rows returned. How is this brand new Hive table already populated with records? _____

3.5. Why did the previous query not require a MapReduce job to execute?

3.6. Try the following query. Make sure the output looks like first names:

```
hive> select fname from wh_visits limit 20;
```

This time a MapReduce job executed. Why? _____

Step 4: Count the Number of Rows in a Table

4.1. Enter the following Hive query, which outputs the number of rows in **wh_visits**:

```
hive> select count(*) from wh_visits;
```

4.2. How many rows are currently in **wh_visits**? _____

Step 5: Selecting the Input File Name

5.1. Hive has two virtual columns that gets created automatically for every table: **INPUT_FILE_NAME** and **BLOCK_OFFSET_INSIDE_FILE**. You can use these column names in your queries just like any other column of the table. To demonstrate, run the following query:

```
hive> select INPUT_FILE_NAME, lname, fname FROM wh_visits
WHERE lname LIKE 'Y%';
```

5.2. The result of this query is visitors to the White House whose last name starts with "Y". Notice that the output also contains the particular file that the record was found in:

```
hdfs://sandbox.hortonworks.com:8020/apps/hive/warehouse/wh_visits/part-m-00001 YOUNGLEDISI
hdfs://sandbox.hortonworks.com:8020/apps/hive/warehouse/wh_visits/part-m-00002 YARNOLD DAVID
```

Step 6: Drop a Table

6.1. Let's see what happens when a managed table is dropped. Start by defining a simple table called **names** using the Hive Shell:

```
hive> create table names (id int, name string)
> ROW FORMAT DELIMITED FIELDS TERMINATED BY '\t';
```

6.2. Use the Hive **dfs** command to put **Lab7.1/names.txt** into the table's warehouse folder:

```
hive> dfs -put /root/labs/Lab7.1/names.txt
/apps/hive/warehouse/names/;
```

6.3. View the contents of the table's warehouse folder:

```
hive> dfs -ls /apps/hive/warehouse/names;  
Found 1 items  
root hdfs    78 /apps/hive/warehouse/names/names.txt
```

6.4. From the Hive Shell, run the following query:

```
hive> select * from names  
; OK  
0    Rich  
1    Barry  
2    George  
3    Ulf  
4    Danielle  
5    Tom  
6    manish  
7    Brian  
8    Mark
```

6.5. Now drop the **names** table:

```
hive> drop table names;
```

6.6. View the contents of the table's warehouse folder again. Notice the **names** folder is gone:

```
hive> dfs -ls /apps/hive/warehouse/names;  
ls: '/apps/hive/warehouse/names': No such file or directory
```

IMPORTANT: Be careful when you drop a managed table in Hive. Make sure you either have the data backed up somewhere else, or that you no longer want the data.

Step 7: Define an External Table

7.1. In this step you will see how external tables work in Hive. Start by putting **names.txt** into HDFS:

```
hive> dfs -put /root/labs/Lab7.1/names.txt names.txt;
```

7.2. Create a folder in HDFS for the external table to store its data in:

```
hive> dfs -mkdir hivedemo;
```

7.3. Define the **names** table as external this time:

```
hive> create external table names (id int, name string)
> ROW FORMAT DELIMITED FIELDS TERMINATED BY '\t'
> LOCATION '/user/root/hivedemo';
```

7.4. Load data into the table:

```
hive> load data inpath '/user/root/names.txt' into table names;
```

7.5. Verify the load worked:

```
hive> select * from names;
```

7.6. Notice the **names.txt** file has been moved to **/user/root/hivedemo**:

```
hive> dfs -ls hivedemo; Found 1 items
-rw-r--r--  1 root hdfs    78 hivedemo/names.txt
```

7.7. Similarly, verify **names.txt** is no longer in your **/user/root** folder in HDFS. Why is it gone?

7.8. Use the **ls** command to verify that the **/apps/hive/warehouse** folder does not contain a subfolder for the **names** table.

7.9. Now drop the **names** table:

```
hive> drop table names;
```

7.10. View the contents of **/user/root/hivedemo**. Notice that **names.txt** is still there.

RESULT: As you just verified, the data for external tables is not deleted when the corresponding table is dropped. Aside from this behavior, managed tables and external tables in Hive are essentially the same. You now have a table in Hive named **wh_visits** that was loaded from the result of a Pig job. You also have an external table called **names** that stores its data in **/user/root/hivedemo**. At this point, you should have a good understanding of how Hive tables are created and populated.

Lab 2: Analyzing Big Data with Hive

Perform the following steps:

Step 1:

1.1. Create a new text file named `whitehouse.hive` and save it in your Lab7.2 folder.

1.2. In this step, you will find the first visitor to the White House (based on our dataset). This will involve some clever handling of timestamps. This will be a long query, so enter it on multiple lines. Start by selecting all columns where the **time_of_arrival** is not empty:

```
select * from wh_visits where time_of_arrival != ""
```

1.3. To find the first visit, we need to sort the result. This requires converting the **time_of_arrival** string into a timestamp. We will use the **unix_timestamp** function to accomplish this. Add the following **order by** clause:

```
order by unix_timestamp(time_of_arrival,  
'MM/dd/yyyy hh:mm')
```

1.4. Since we are only looking for one result, we certainly don't need to return every row. Let's limit the result to 10 rows, so we can view the first 10 visitors:

```
limit 10;
```

1.5. Save your changes to **whitehouse.hive**.

1.6. Execute the script **whitehouse.hive** and wait for the results to be displayed:

```
# hive -f whitehouse.hive
```

1.7. The results should be 10 visitors, and the first visit should be in 2009 since that is when the dataset begins. The first visitor is Charles Kahn on 3/5/2009.

Step 2: Find the Last Visit

2.1. This one is easy – just take the previous query and reverse the order by adding **desc** to the **order by** clause:

```
order by unix_timestamp(time_of_arrival,  
'MM/dd/yyyy hh:mm') desc
```

2.2. Run the query again, and you should see that the most recent visit was Jackie Walker on 3/18/2011.

Step 3: Find the Most Common Comment

3.1. In this step, you will explore the **info_comment** field and try to determine the most common comment. You will use some of Hive's aggregate functions to accomplish this. Start by creating a new text file named **comments.hive**.

3.2. You will now create a query that displays the 10 most frequently occurring comments. Start with the following select clause:

```
from wh_visits  
select count(*) as comment_count, info_comment
```

This runs the aggregate **count** function on each group (which you will define later in the query) and names the result **comment_count**. For example, if "OPEN HOUSE" occurs 5 times, then **comment_count** will be 5 for that group.

Notice we are also selecting the **info_comment** column so we can see what the comment is.

3.3. Group the results of the query by the **info_comment** column:

```
group by info_comment
```

3.4. Order the results by **comment_count**, because we are only interested in comments that appear most frequently:

```
order by comment_count DESC
```

3.5. We only want the top results, so limit the result set to 10:

```
limit 10;
```

3.6. Save your changes to **comments.hive** and execute the script. Wait for the MapReduce job to execute.

3.7. The output will be 10 comments and should look like:

```
9036
1253 HOLIDAY BALL ATTENDEES/
894  WHO EOP RECEP 2
700  WHO EOP 1 RECEPTION/
601  RESIDENCE STAFF HOLIDAY RECEPTION/
586  PRESS RECEPTION ONE (1)/
580  GENERAL RECEPTION 1
540  HANUKKAH RECEPTION./
540  GEN RECEP 5/
```

3.8. It appears that a blank comment is the most frequent comment, followed by the HOLIDAY BALL, then a variation of other receptions.

3.9. OPTIONAL: Modify the query so that it ignores empty comments. If it works, the comment "GEN RECEP 6/" will show up in your output.

Step 4: Least Frequent Comment

4.1. Run the previous query again, but this time find the 10 least occurring comments. The output should look like:

```
1 merged to u59031
1 WHO EOP/
1 WHO EOP RECLEAR
1 WAITING FOR SUPERMAN VISIT
1 ST. PATRICK'S RECEPTION GUESTS
1 SCIENCE FAIR
1 RES PARTY/
1 PRIVATE MEETING
1 PRIVATE LUNCH
1 POTUS PHOTO W/ US ATTORNEYS/
```

Hive Programming Guide

This seems accurate since 1 is the least number of times a comment can appear. Plus this query reveals that Superman has visited the President at least once!

Step 5: Analyze the Data Inconsistencies

5.1. Analyzing the results of the most and least frequent comments, it appears that several variations of GENERAL RECEPTION occur. In this step, you will try to determine the number of visits to the POTUS involving a general reception by trying to clean up some of these inconsistencies in the data.

NOTE: Inconsistencies like these are very common in big data, especially when human input is involved. In this dataset, we like have different people entering similar comments but using their own abbreviations.

5.2. Modify the query in **comments.hive**. Instead of searching for empty comments, search for comments that contain the string "RECEP".

```
where info_comment like "%RECEP%"
```

5.3. Change the limit clause from 10 to 30:

```
limit 30;
```

5.4. Run the query again.

5.5. Notice there are several GENERAL RECEPTION entries that only differ by a number at the end, or use the GEN RECEP abbreviation:

```
580 GENERAL RECEPTION 1
540 GEN RECEP 5/
516 GENERAL RECEPTION 3
498 GEN RECEP 6/
438 GEN RECEP 4
31 GENERAL RECEPTION 2
23 GENERAL RECEPTION 3
20 GENERAL RECEPTION 6
20 GENERAL RECEPTION 5
13 GENERAL RECEPTION 1
```

5.6. Let's try one more query to try and narrow GENERAL RECEPTION visits. Modify the WHERE clause in comments.hive to include "%GEN%": where info_comment like "%RECEP%" and info_comment like "%GEN%"

5.7. Leave the limit at 30, and run the query again.

5.8. The output this time reveals all the variations of GEN and RECEP. Let's add up the total number of them by running the following query:

```
from wh_visits
select count(*)
where info_comment like "%RECEP%" and info_comment like "%GEN%";
```

5.9. Notice there are 2,697 visits to the POTUS with GEN RECEP in the comment field, which is about 12% of the 21,819 total visits to the POTUS in our dataset.

NOTE: More importantly, these results show that our first query of 1,253 attendees to the HOLIDAY BALL does not mean that the holiday ball is the most likely reason to visit the President. More than twice as many visitors are there for a general reception. This type of analysis is common in big data, and it shows how Data Analysts need to be creative when researching their data.

Step 6: Verify the Result

6.1. We have 12% of visitors to the POTUS going for a general reception, but there were a lot of statements in the comments that contained WHO and EOP. Modify the query from the last step and display the top 30 comments that contain "WHO" and "EOP". The result should look like:

```
894 WHO EOP RECEP 2
700 WHO EOP 1 RECEPTION/
43 WHO EOP RECEP/
20 WHO EOP HOLIDAY RECEP/
13 WHO/EOP #2/
8 WHO EOP RECEPTION
```

```
7   WHO EOP RECEP
1   WHO EOP/
1   WHO EOP RECLEAR
```

6.2. Run a query that counts the number of records with WHO and EOP in the comments:

```
from wh_visits select count(*) where info_comment like "%WHO%" and
info_comment like "%EOP%";
```

You should get 1,687 visits, or 7.7% of the visitors to the POTUS. So GENERAL RECEPTION still appears to be the most frequent comment.

Step 7: Find the Most Visits

7.1. See if you can write a Hive script that finds the top 20 individuals who visited the POTUS most. Use the Hive command from Step 3 earlier in this lab as a guide. **HINT:** Use a grouping by both **fname** and **lname**.

7.2. To verify your script worked, here are the top 20 individuals who visited the POTUS, along with the number of visits:

```
16  ALAN PRATHER
15  CHRISTOPHER FRANKE
15  ANNAMARIA MOTTOLA
14  ROBERT BOGUSLAW
14  CHARLES POWERS
12  SARAHHART
12  JACKIE WALKER
12  IASONFETTIG
12  SHENGTSUNG WANG
12  FERN SATO
12  DIANAFISH
11  JANETBAILEY
11  PETERWILSON
11  GLENNDEWEY
11  MARCIO BOTELHO
11  DONNAWILLINGHAM
10  DAVIDAXELROD
10  CLAUDIA CHUDACOFF
10  VALERIE JARRETT
10  MICHAEL COLBURN
```

Hive Programming Guide

RESULT: You have written several Hive queries to analyze the White House visitor data. The goal is for you to become comfortable with working with Hive, so hopefully you now feel like you can tackle a Hive problem and be able to answer questions about your big data stored in Hive.

Lab 3: Understanding MapReduce in Hive

Step 1: The Describe Command

1.1. Run the `describe` command on the **wh_visits** table:

```
hive> describe wh_visits;
OK
lname          string      None
fname          string      None
time_of_arrival string      None
appt_scheduled_time string     None
meeting_location string     None
info_comment    string      None
Time taken: 0.677 seconds
```

1.2. Did this query require a MapReduce job? _____

1.3. What is the name of the Hive resource that was accessed to retrieve this schema information? _____

Step 2: A Simple Query

2.1. Run the following query:

```
select * from wh_visits where fname = "JOE";
```

2.2. Does Hive run a MapReduce job to generate the result? _____

2.3. Open your browser and point it to the JobHistory UI:

```
http://ipaddress:19888/
```

2.4. Notice the most recent MapReduce job executed is your "JOE" query:

2.5. How many map tasks were used to execute this query? _____

2.6. How many reduce tasks were used to execute this query? _____

2.7. How many attempts did it take for this task to succeed? _____

2.8. How long did it take for this query to execute? _____

Step 3: A Sorted Query

3.1. Run the following query:

```
hive> select * from wh_visits where fname = "JOE" sort by lname;
```

3.2. When the MapReduce job completes, find its job details page from the Job Browser.

3.3. How many map tasks were used to execute this query? _____

3.4. How many reduce tasks were used to execute this query? _____

3.5. The map task outputs <key,value> pairs and sends them to the reducer. What do you think this MapReduce job chose as the key for the mapper's output? _____

Step 4: A Select Query

4.1. Run the following query:

```
hive> select * from wh_visits limit 5;
```

4.2. Does Hive run a MapReduce job to generate the results? _____

4.3. What data is read from HDFS? _____

4.4. Now select a single column from **wh_visits**:

```
hive> select fname from wh_visits limit 5;
```

4.5. Why did this require a MapReduce job but “**select ***” did not? _____

Step 5: Using the **EXPLAIN** Command

5.1. The **EXPLAIN** command shows the execution plan of a query, without actually executing the query. To demonstrate, add **EXPLAIN** to the beginning of the following query that you ran earlier in this lab:

```
hive> explain select * from wh_visits where fname = "JOE" sort by lname;
```

5.2. Notice the query is executed in two stages. **Stage-0** performs the limit operator. This was the first mapper that executed the query.

5.3. Notice **Stage-1** is a MapReduce job that has one mapper (look for **Map Operator Tree**) and one reducer (under **Reduce Operator Tree**). As you can see from this execution plan, the mapper is doing most of the work.

Step 6: Use **EXPLAIN EXTENDED**

6.1. Run the previous **EXPLAIN** again, except this time add the **EXTENDED** command:

```
hive> explain extended select * from wh_visits where fname  
= "JOE" sort by lname;
```

6.2. Compare the two outputs. Notice the **EXTENDED** command adds a lot of additional information about the underlying execution plan.

ANSWERS:

1.2: No

1.3: The Hive metastore contains the schema information of all tables.

2.3: Yes

2.5: 1 map task

2.6: 0 reduce tasks

2.7: It probably succeeded on the first attempt. If not, you will see multiple entries in the list of attempts.

2.8: Subtract the Execution Finish Time from the Execution Start Time. It should have executed in about 15-30 seconds.

3.3: 1 map task

3.4: 1 reduce task

3.5: It makes sense for the mapper to use **lname** as the key, which would mean the visitors would already be sorted by last name when they got to the reducer.

4.2: No

4.3: Hive simply reads the data directly from the underlying file in HDFS.

4.5: Selecting specific columns requires actual processing of the contents of each record to split out the required columns.

Lab 4: Joining Datasets in Hive

Location of Files:	/root/labs/Lab7.4
--------------------	-------------------

Perform the following steps:

Step 1: Load the Data into Hive

1.1. View the contents of the file **setup.hive** in **/root/labs/Lab7.4**:

```
# more setup.hive
```

1.2. Notice this script creates three tables in Hive. The **nyse_data** table is filled with the daily stock prices of stocks that start with the letter “K”, and the **dividends** table that contains the quarterly dividends of those stocks. The **stock_aggregates** table is going to be used for a join of these two datasets and contain the stock price and dividend amount on the date the dividend was paid.

1.3. Run the **setup.hive** script from the **Lab7.4** folder:

```
# hive -f setup.hive
```

1.4. To verify the script worked, enter the following query from the Hive Shell:

```
hive> select * from nyse_data limit 20;
hive> select * from dividends limit 20
```

You should see daily stock prices and dividends from stocks that start with the letter “K”.

1.5. The **stock_aggregates** table should be empty, but view its schema to verify it was created successfully:

```
hive> describe stock_aggregates;
```

```
OK
symbol          string      None
year            string      None
high            float       None
low             float       None
average_close    float       None
total_dividends  float       None
```

Step 2: Join the Datasets

2.1. The **join** statement is going to be fairly long, so let's create it in a text file. Create a new text in the **Lab7.4** folder named **join.hive**, and open the file with a text editor.

2.2. We will break the join statement down into sections. First, the result of the **join** is going to put into the **stock_aggregates** table, which requires an **insert**:

```
insert overwrite table stock_aggregates
```

The **overwrite** causes any existing data in **stock_aggregates** to be deleted.

2.3. The data being inserted is going to be the result of a select query that contains various insightful indicators about each stock. The result is going to contain the stock symbol, date traded, maximum high for the stock, minimum low, average close, and sum of dividends, as shown here:

```
select a.symbol, year(a.trade_date), max(a.high), min(a.low), avg(a.close),
sum(b.dividend)
```

2.4. The from clause is the **nyse_data** table:

```
from nyse_data a
```

2.5. The join is going to be a left outer join of the **dividends** table:

```
left outer join dividends b
```

2.6. The join is by stock symbol and trade date:

```
on (a.symbol = b.symbol and a.trade_date = b.trade_date)
```

2.7. Let's group the result by symbol and trade date:

```
group by a.symbol, year(a.trade_date);
```

2.8. Save your changes to join.hive.

Step 3: Run the Query

3.1. Run the query and wait for the MapReduce jobs to execute:

```
# hive -f join.hive
```

3.2. How many MapReduce jobs does it take to perform this query?

Step 4: Verify the Results

4.1. Run a **select** query to view the contents of **stock_aggregates**:

```
hive> select * from stock_aggregates;
```

The output should look like:

```
KYO 2004 90.9 66.2575.79952 0.544
KYO 2005 78.4562.5872.042656 0.91999996
KYO 2006 98.0171.7385.80327 0.851
KYO 2007 110.01 81.0 93.737686 NULL
KYO 2008 100.78 45.4179.6098 NULL
KYO 2009 93.2 52.9877.04389 NULL
KYO 2010 93.8385.9490.71NULL
stock_symbol NULL NULL NULL NULL NULL
```

4.2. List the contents of the **stock_aggregates** directory in HDFS. The **000000_0** file was created as a result of the **join** query:

```
# hadoop fs -ls -R /apps/hive/warehouse/stock_aggregates/
41109 /apps/hive/warehouse/stock_aggregates/000000_0
```

4.3. View the contents of the **stock_aggregates** table using the **cat** command:

```
# hadoop fs -cat
/apps/hive/warehouse/stock_aggregates/000000_0
```

RESULT: The **stock_aggregates** table is a joining of the daily stock prices and the quarterly dividend amounts on the date the dividend was announced, and the data in table is an aggregate of various statistics like max high, min low, etc.

Lab 5: Computing ngrams of Emails in Avro Format

Location of Files:	/root/labs/Lab7.5
--------------------	-------------------

Step 1: View an Avro Schema

1.1. Change directories to the **Lab7.5** folder. Notice this folder contains an Avro file named **sample.avro**.

1.2. Enter the following command to view the schema of the contents of **sample.avro**:

```
avro cat --print-schema sample.avro
```

1.3. How many fields do records in **sample.avro** have? _____

1.4. Create a schema file for **sample.avro**:

```
avro cat --print-schema sample.avro > /tmp/sample.avsc
```

Step 2: Create a Hive Table from an Avro Schema

2.1. View the contents of the **CREATE TABLE** query defined in the **create_sample_table.hive** file in your **Lab7.5** folder.

2.2. Make sure the **avro.schema.file** property points to the schema file you created in the previous step:

```
WITH SERDEPROPERTIES (  
'avro.schema.url'='file:///tmp/sample.avsc')
```

2.3. Run the **CREATE TABLE** query:

```
hive -f create_sample_table.hive
```

Step 3: Verify the Table

3.1. Start the Hive shell.

3.2. Run the **show tables** command and verify that you have a table named **sample_table**.

3.3. Run the describe command on **sample_table**. Notice the schema for **sample_table** matches the Avro schema from **sample.avsc**.

3.4. Let's associate some data with **sample_table**. Copy **sample.avro** into the Hive **warehouse** folder by running the following command (all on a single line):

```
hive> dfs -put /root/labs/Lab7.5/sample.avro  
/apps/hive/warehouse/sample_table
```

3.5. View the contents of **sample_table**:

```
hive> select * from sample_table;  
OK  
Foo 19 10, Bar Eggs Spam 800
```

Note there is only one record in **sample.avro**.

Step 4: Create Email User Table

4.1. There is an Avro file in your **Lab7.5** folder named **mbox7.avro**, which represents emails in an Avro format from a Hive mailing list for the month of July. Use the ~~---print-schema~~ option of **avro** to view the schema of this file.

4.2. How many fields do records in **mbox7.avro** have? _____

4.3. Run the ~~---print-schema~~ command again, but this time output the schema to a file named **mbox.avsc**:

```
avro cat --print-schema mbox7.avro > /tmp/mbox.avsc
```

4.4. View the contents of the **create_email_table.hive** script in your **Lab7.5** folder. Verify the **avro.schema.url** property is correct.

4.5. Run the script to create the **hive_user_email** table:

```
hive -f create_email_table.hive
```

4.6. Copy **mbox7.avro** into the warehouse directory:

```
hadoop fs -put mbox7.avro  
/apps/hive/warehouse/hive_user_email/
```

4.7. Verify the table has data in it:

```
select * from hive_user_email limit 20;
```

Step 5: Compute a Bigram

5.1. Start the Hive shell.

5.2. Use the Hive **ngrams** function to create a bigram of the words in **mbox7.avro**:

```
select  
ngrams(sentences(content),2 ,10)  
from hive_user_email;
```

The output will be kind of a jumbled mess:

```
[{"ngram":["2013","at"],"estfrequency":802.0}, {"ngram":["of","the"],"estfrequency":391.0}, {"ngram":["I","am"],"estfrequency":368.0}, {"ngram":["I","have"],"estfrequency":340.0}, {"ngram":["J","E9r"],"estfrequency":306.0}, {"ngram":["for","the"],"estfrequency":291.0}, {"ngram":["you","are"],"estfrequency":289.0}, {"ngram":["user","hive.apache.org"],"estfrequency":289.0}, {"ngram":["to","the"],"estfrequency":276.0}, {"ngram":["E9r","F4me"],"estfrequency":270.0}]
```

5.3. To clean this up, use the Hive **explode** function to display the output in a more readable format:

```
select
explode(ngrams(sentences(content),2 ,10)) as x from hive_user_email;
```

You should see the a nice, readable list of 10 bigrams:

```
{"ngram":["2013","at"],"estfrequency":802.0}
{"ngram":["of","the"],"estfrequency":391.0}
{"ngram":["I","am"],"estfrequency":368.0}
{"ngram":["I","have"],"estfrequency":340.0}
{"ngram":["J","E9r"],"estfrequency":306.0}
{"ngram":["for","the"],"estfrequency":291.0}
{"ngram":["you","are"],"estfrequency":289.0}
{"ngram":["user","hive.apache.org"],"estfrequency":289.0}
{"ngram":["to","the"],"estfrequency":276.0}
{"ngram":["E9r","F4me"],"estfrequency":270.0}
```

5.4. Typically when working with word comparison we ignore case. Run the query again, but this time add the Hive **lower** function and compute 20 bigrams:

```
select
explode(ngrams(sentences(lower(content)),2 ,20)) as x from
hive_user_email;
```

The output should look like the following:

```
{"ngram":["2013","at"],"estfrequency":802.0}
{"ngram":["i","have"],"estfrequency":409.0}
{"ngram":["of","the"],"estfrequency":391.0}
{"ngram":["i","am"],"estfrequency":372.0}
{"ngram":["if","you"],"estfrequency":347.0}
{"ngram":["in","hive"],"estfrequency":337.0}
{"ngram":["for","the"],"estfrequency":309.0}
{"ngram":["j","e9r"],"estfrequency":306.0}
{"ngram":["you","are"],"estfrequency":289.0}
{"ngram":["user","hive.apache.org"],"estfrequency":289.0}
{"ngram":["to","the"],"estfrequency":276.0}
{"ngram":["outer","join"],"estfrequency":271.0}
{"ngram":["2013","06"],"estfrequency":270.0}
{"ngram":["e9r","f4me"],"estfrequency":270.0}
{"ngram":["left","outer"],"estfrequency":270.0}
{"ngram":["in","the"],"estfrequency":252.0}
{"ngram":["gmail.com","wrote"],"estfrequency":248.0}
{"ngram":["17","16"],"estfrequency":248.0}
{"ngram":["06","17"],"estfrequency":246.0}
{"ngram":["wrote","hi"],"estfrequency":234.0}
```

Step 6: Compute a Context n-gram

6.1. From the Hive shell, run the following query, which uses the `context_ngrams` function to find the top 20 terms that follow the word **“error”**:

```
select explode(context_ngrams(sentences(lower(content)),
array("error", null),20)) as x
from hive_user_email;
```

The output should look like the following:

```
{"ngram":["in"],"estfrequency":102.0}
{"ngram":["return"],"estfrequency":97.0}
{"ngram":["org.apache.hadoop.hive ql.exec.udfargumenttypeex
ception"],"estfrequency":49.0}
{"ngram":["failed"],"estfrequency":49.0}
{"ngram":["is"],"estfrequency":41.0}
{"ngram":["message"],"estfrequency":40.0}
{"ngram":["when"],"estfrequency":39.0}
{"ngram":["please"],"estfrequency":36.0}
{"ngram":["while"],"estfrequency":28.0}
{"ngram":["org.apache.thrift.transport.ttransportexception"
],"estfrequency":28.0}
{"ngram":["datanucleus.plugin"],"estfrequency":26.0}
{"ngram":["during"],"estfrequency":18.0}
{"ngram":["query"],"estfrequency":16.0}
{"ngram":["hive"],"estfrequency":16.0}
{"ngram":["could"],"estfrequency":16.0}
{"ngram":["java.lang.runtimeexception"],"estfrequency":13.0
}
{"ngram":["13"],"estfrequency":12.0}
{"ngram":["error"],"estfrequency":12.0}
{"ngram":["exec.execdriver"],"estfrequency":10.0}
{"ngram":["exec.task"],"estfrequency":10.0}
```

6.2. What is the most likely word to follow “**error**” in these emails? _____

6.3. Run a Hive query that finds the top 20 results for words in mbox7.avro that follow the phrase “error in”.

RESULT: You have written several Hive queries that computed bigrams based on the data in the mbox7.avro file, and also computed histograms using the data in the orders table. You should also be familiar with working with Avro files, a popular file format in Hadoop.

SOLUTIONS:

Step 6.3:

```
select explode(context_ngrams(sentences(lower(content)),  
array("error", "in", null) ,20)) as x from hive_user_email;
```

Lab 7: Using HCatalog with Pig

Location of Files:	n/a
--------------------	-----

Perform the following steps:

Step 1: Start the Grunt Shell

1.1. SSH into your HDP 2.0 virtual machine.

1.2. Start the Grunt shell for use with HCatalog:

```
# pig -useHCatalog
```

Step 2: Load a HCatalog Table

2.1. Define a relation for the **wh_visits** table in Hive using the **HCatLoader()**:

```
grunt> visits = LOAD 'wh_visits' USING  
org.apache.hcatalog.pig.HCatLoader();
```

2.2. View the schema of the **visits** relation to verify it matches the schema of the **wh_visits** table:

```
visits: {lname: chararray,fname: chararray,time_of_arrival:  
chararray,appt_scheduled_time: chararray,meeting_location:  
chararray,info_comment: chararray}
```

Step 3: Run a Pig Query

3.1. Let's execute a query to verify everything is working. Define the following relation:

```
grunt> joe = FILTER visits BY (fname == 'JOE');
```


3.2. Dump the relation:

```
grunt> DUMP joe;
```

The output should be visitors from **wh_visits** with the firstname "JOE".

Step 4: Create an HCatalog Schema

4.1. Quit the Grunt shell and start the Hive shell.

4.2. An HCatalog schema is essentially just a table in the Hive metastore. To define a schema for use with HCatalog, create a table in Hive:

```
hive> create table joes (fname string, lname string, comments string);
```

4.3. Verify the table was created successfully using '*show tables*'.

4.4. Use the **describe** command to view the schema of **joes**:

```
hive> describe joes;
OK
fname                string                None
lname                string                None
comments             string                None
```

Step 5: Using HCatStorer

5.1. Exit the Hive shell and start the Grunt shell again. Be sure to use the useHCatalog option:

```
# pig -useHCatalog
```

5.2. Define the **visits** and **joe** relations again (using the up arrow to browse through the history of Pig commands).

5.3. In this step, you will use **HCatStorer** in Pig to input records into the **joes** table. To do this, you need a relation whose fields match the schema of **joes**. You can accomplish this using a projection. Define the following relation:

```
grunt> project_joe = FOREACH joe GENERATE fname, lname, info_comment;
```

5.4. Store the projection into the HCatalog table using the **STORE** command:

```
grunt> STORE project_joe INTO 'joes' USING  
org.apache.hcatalog.pig.HCatStorer();
```

This command failed. Why? _____

5.5. Notice the projection has fields named **fname**, **lname** and **info_comment**, but the **joes** table in HCatalog has a schema with **fname**, **lname** and **comments**. The **fname** and **lname** fields match, but **info_comment** needs to be renamed to **comments**. Modify your projection by using the **AS** keyword:

```
project_joe = FOREACH joe GENERATE fname, lname, info_comment AS  
comments;
```

5.6. Now run the **STORE** command again:

```
grunt> STORE project_joe INTO 'joes' USING  
org.apache.hcatalog.pig.HCatStorer();
```

This time the command should work and a MapReduce job will execute.

Step 6: Verify the STORE Worked

6.1. Quit the Grunt shell and start the Hive shell again.

6.2. View the contents of the **joes** table:

```
hive> select * from joes;
```

You should see visitors all named “JOE”, along with their last name and the comments.

Step 7: View the Files

7.1. You can also check the file system to see if a **STORE** command worked. From the command line, view the contents of **/apps/hive/warehouse/joes**:

```
# hadoop fs -ls /apps/hive/warehouse/joes/ Found 1 items
root hdfs      896 /apps/hive/warehouse/joes/part-m-00000
```

Notice that the file for the **joes** table is named **part-m-00000**. Where did that name come from? _____

7.2. Use the **cat** command to view the contents of **part-m-00000**:

```
# hadoop fs -cat /apps/hive/warehouse/joes/part-m-00000
```

As you can see, this is the same list of names from the Hive select * query, which should be no surprise at this point in the course!

RESULT: You have seen how to run a Pig script that uses HCatalog to provide the schema using HCatLoader and HCatStorer.

ANSWERS:

5.4: The initial **STORE** command failed because the field names in the relation you were trying to store did not match the column names of the underlying table's schema.

7.1: The part---m--00000 file is a result of the Pig MapReduce job that executed when you ran the STORE command with HCatStorer.

Lab 8 : Advanced Hive Programming

Location of Files:	/root/labs/Lab9.1
--------------------	-------------------

Step 1: Create and Populate a Hive Table

1.1. From the command line, change directories to the **Lab9.1** folder:

```
# cd ~/labs/Lab9.1
```

1.2. View the contents of the **orders.hive** file in the **Lab9.1** folder:

```
# more orders.hive
```

Notice it defines a Hive table named orders that has 7 columns.
Notice it also loads the contents of **/tmp/shop.tsv** into the orders table.

1.3. Copy **shop.tsv** into the **/tmp** folder:

```
# cp shop.tsv /tmp/
```

1.4. Execute the contents of **orders.hive**:

```
# hive -f orders.hive
```

1.5. From the Hive shell, verify the script worked by running the following two commands:

```
hive> describe orders;  
hive> select count(*) from orders;
```

Your **orders** table should contain 99,999 records.

Step 2: Analyze the Customer Data

2.1. Let's run a few queries to see what this data looks like. Start by verifying that the **username** column actually looks like names:

```
hive> SELECT username FROM orders LIMIT 10;
```

You should see ten first names.

2.2. The orders table contains orders placed by customers. Run the following query, that shows the 10 lowest-price orders:

```
hive> SELECT username, ordertotal FROM orders ORDER BY  
ordertotal LIMIT 10;
```

The smallest orders are each \$10, as you can see from the output:

```
Jeremy10  
Christina 10  
Jasmine 10  
Hannah10  
Thomas10  
Michelle 10  
Brian 10  
Amber 10  
Maria 10  
Victoria 10
```

2.3. Run the same query, but this time use descending order:

```
hive> SELECT username, ordertotal FROM orders ORDER BY  
ordertotal DESC LIMIT 10;
```

The output this time is the 10 highest---priced orders:

```
Mark 612
Jordan612
Anthony 612
Brandon 612
Sean 612
Paul 611
Nathan611
Eric 611
Jonathan 611
Andrew610
```

2.4. Let's find out if men or women spent more money:

```
hive> SELECT sum(ordertotal), gender
FROM orders GROUP BY gender;
```

Based on the output, which gender has spent more money on purchases? _____

2.5. The **orderdate** column is a string with the format **yyyy-mm-dd**. Use the **year** function to extract the various parts of the date. For example, run the following query, which computes the sum of all orders for each year:

```
SELECT sum(ordertotal), year(order_date) FROM orders GROUP BY
year(order_date);
```

The output should look like:

```
4082780 2009
4404806 2010
4399886 2011
4248950 2012
2570749 2013
```

Step 3: Multi-File Insert

3.1. In this step, you will run two completely different queries, but in a single MapReduce job. The output of the queries will be in two separate directories in HDFS. Start by creating a new text file in the **Lab9.1** folder named **multifile.hive**.

3.2. Within the text file, enter the following query. Notice there is no semi-colon between the two **INSERT** statements:

```
FROM ORDERS o
  INSERT OVERWRITE DIRECTORY '2010_orders' SELECT o.* WHERE
year(order_date) = 2010
  INSERT OVERWRITE DIRECTORY 'software'
  SELECT o.* WHERE itemlist LIKE '%Software%';
```

3.3. Save your changes to **multifile.hive**.

3.4. Run the query from the command line:

```
# hive -f multifile.hive
```

3.5. The above query executes in a single MapReduce job. Even more interesting, it only requires a map phase. Why did this job not require a reduce phase?

3.6. Verify the two queries executed successfully by viewing the folders in HDFS:

```
# hadoop fs -ls
```

You should see two new folders: **2010_orders** and **software**.

3.7. View the output files in these two folders. Verify the **2010_orders** directory contains orders from only the year 2010, and verify the **software** directory contains only orders that included '**Software**'.

Step 4: Define a View

4.1. Define a view named **2013_orders** that contains the **orderid**, **order_date**, **username**, and **itemlist** columns of the **orders** table where the **order_date** was in the year 2013.

4.2. Run the **show tables** command:

```
hive> show tables;
```

You should see **2013_orders** in the list of tables.

4.3. To verify your view is defined correctly, run the following query:

```
hive> SELECT COUNT(*) FROM 2013_orders;
```

The **2013_orders** view should contain 13,104 records.

Step 5: Find the Maximum Order of Each Customer

5.1. Suppose you want to find the maximum order of each customer. This can be done easily enough with the following Hive query. Run this query now:

```
hive> SELECT max(ordertotal), userid
```

```
FROM orders GROUP BY userid;
```

5.2. How many different customers are in the **orders** table? _____

5.3. Suppose you want to add the **itemlist** column to the previous query. Try adding it to the **SELECT** clause:

```
hive> SELECT max(ordertotal), userid, itemlist  
FROM orders GROUP BY userid;
```

Notice this query is not valid because **itemlist** is not in the **GROUP BY** key.

5.4. We can join the result set of the max-total query with the **orders** table to add the **itemlist** to our result. Start by defining a view named **max_ordertotal** for the maximum order of each customer:

```
hive> CREATE VIEW max_ordertotal AS
SELECT max(ordertotal) AS maxtotal, userid
FROM orders GROUP BY userid;
```

5.5. Now join the **orders** table with your **max_ordertotal** view:

```
hive> SELECT ordertotal, orders.userid, itemlist
FROM orders
JOIN max_ordertotal ON max_ordertotal.userid = orders.userid AND
max_ordertotal.maxtotal = orders.ordertotal;
```

5.6. How many MapReduce jobs did this query need? _____

5.7. The end of your output should look like:

```
600 98
Grill,Freezer,Bedding,Headphones,DVD,Table,Grill,Software,D
ishwasher,DVD,Microwave,Adapter
600 99 Washer,Cookware,Vacuum,Freezer,2-Way
Radio,Bicycle,Washer & Dryer,Coffee
Maker,Refrigerator,DVD,Boots,DVD
600 100 Bicycle,Washer,DVD,Wrench Set,Sweater,2-Way
Radio,Pants,Freezer,Blankets,Grill,Adapter,pillows
```

NOTE: In the next lab, you will optimize this query using a custom Python script to avoid the need for two MapReduce jobs.

Step 6: Fixing the GROUP BY Key Error

6.1. Let's compute the sum of all of the orders of all customers. Start by entering the following query:

```
SELECT sum(ordertotal), userid FROM orders GROUP BY userid;
```

Notice the output is the sum of all orders, but displaying just the userid is not very exciting.

6.2. Try to add the **username** column to the **SELECT** clause:

```
SELECT sum(ordertotal), userid, username  
FROM orders  
GROUP BY userid;
```

This generates the infamous “Expression not in GROUP BY key” error, because **username** column is not being aggregated but the **ordertotal** is.

6.3. An easy fix is to aggregate the **username** values using the **collect_set** function, but output only one of them:

```
SELECT sum(ordertotal), userid, collect_set(username)[0] FROM orders  
GROUP BY userid;
```

You should get the same output as before, but this time the **username** is included.

Step 7: Using the **OVER** Clause

7.1. Now let's compute the sum of all orders for each customer, but this time use the **OVER** clause to not group the output and to also display the **itemlist** column:

```
SELECT userid, itemlist, sum(ordertotal) OVER (PARTITION BY userid)  
FROM orders;
```

Notice the output contains every order, along with the items they purchased and the sum of all the orders ever placed from that particular customer.

Step 8: Using the Window Functions

8.1. It is not difficult to compute the sum of all orders for each day using the **GROUP BY** clause:

```
select order_date, sum(ordertotal) FROM orders
GROUP BY order_date;
```

Run the query above and the tail of the output should look like:

```
2013-07-28 18362
2013-07-29 3233
2013-07-30 4468
2013-07-31 4714
```

8.2. Suppose you want to compute the sum for each day that includes each order. This can be done using a window that sums all previous orders along with the current row:

```
SELECT order_date, sum(ordertotal) OVER
PARTITION BY order_date ROWS BETWEEN UNBOUNDED PRECEDING AND
CURRENT ROW)
FROM orders;
```

To verify it worked, your tail of your output should look like:

```
2013-07-31 3163
2013-07-31 3415
2013-07-31 3607
2013-07-31 4146
2013-07-31 4470
2013-07-31 4610
2013-07-31 4714
```

Step 9: Using the Hive Analytics Functions

9.1. Run the following query, which displays the rank of the **ordertotal** by day:

```
SELECT order_date, ordertotal, rank() OVER  
(PARTITION BY order_date ORDER BY ordertotal) FROM orders;
```

9.2. To verify it worked, the output of July 31, 2013, should look like:

```
2013-07-31 48 1  
2013-07-31 104 2  
2013-07-31 119 3  
2013-07-31 130 4  
2013-07-31 133 5  
2013-07-31 135 6  
2013-07-31 140 7  
2013-07-31 147 8  
2013-07-31 156 9  
2013-07-31 192 10  
2013-07-31 192 10  
2013-07-31 196 12  
2013-07-31 240 13  
2013-07-31 252 14  
2013-07-31 296 15  
2013-07-31 324 16  
2013-07-31 343 17  
2013-07-31 500 18  
2013-07-31 528 19  
2013-07-31 539 20
```

9.3. As a challenge, see if you can run a query similar to the previous one except compute the rank over months, instead of each day.

Step 10: Histograms

10.1. Run the following Hive query, which uses the **histogram_numeric** function to compute 20 (x,y) pairs of the frequency distribution of the total order amount from customers who purchased a microwave (using the **orders** table):

```
select
explode(histogram_numeric(ordertotal,20)) as x from orders
where itemlist LIKE "%Microwave%";
```

The output should look like the following:

```
{"x":14.333333333333332,"y":3.0}
{"x":33.87755102040816,"y":441.0}
{"x":62.52577319587637,"y":679.0}
{"x":89.37823834196874,"y":965.0}
{"x":115.1242236024843,"y":1127.0}
{"x":142.6468885672939,"y":1382.0}
{"x":174.07664233576656,"y":1370.0}
{"x":208.06909090909105,"y":1375.0}
{"x":242.55486381322928,"y":1285.0}

{"x":275.8625954198475,"y":1048.0}
{"x":304.71100917431284,"y":872.0}
{"x":333.1514423076924,"y":832.0}
{"x":363.7630208333335,"y":768.0}
{"x":397.51587301587364,"y":756.0}
{"x":430.9072847682117,"y":604.0}
{"x":461.68715083798895,"y":537.0}
{"x":494.1598360655734,"y":488.0}
{"x":528.5816326530613,"y":294.0}
{"x":555.5166666666672,"y":180.0}
{"x":588.7979797979801,"y":198.0}
```

10.2. Write a similar Hive query that computes 10 frequency distribution pairs for the **ordertotal** from the **orders** table where **ordertotal** is greater than \$200. The output should look like:

```
{"x":218.8195174551819,"y":7419.0}
{"x":254.10237580993478,"y":6945.0}
{"x":293.4231618807192,"y":6338.0}
{"x":334.57302573203015,"y":5635.0}
{"x":379.79714934930786,"y":4841.0}
{"x":428.1165628891644,"y":4015.0}
{"x":473.1484734420741,"y":2391.0}
```

```
{"x":511.2576946288467,"y":1657.0}
{"x":549.0106899902812,"y":1029.0}
{"x":589.0761194029857,"y":670.0}
```

RESULT: You should now be comfortable running Hive queries and using some of the more advanced features of Hive like views and the window functions.

ANSWERS:

2.4: Men spent \$9,919,847 and women spent \$9,787,324.

3.5: Because the query only does a **SELECT ***, no reduce phase was needed.

5.2: There are 100 unique customers in the **orders** table.

5.6: The query resulted in two MapReduce jobs.

SOLUTIONS:

Step 4.1: The **2013_orders** view:

```
CREATE VIEW 2013_orders AS
SELECT orderid, order_date, username, itemlist
FROM orders
WHERE year(order_date) = '2013';
```

Step 9.3: The rank query by month:

```
select substr(order_date,0,7), ordertotal, rank() OVER
(PARTITION BY substr(order_date,0,7) ORDER BY ordertotal) FROM
orders;
```

Step 10.2:

```
select
explode(histogram_numeric(ordertotal,10)) as x from orders
where ordertotal > 200;
```

Lab 9 : Streaming Data with Hive and Python

Location of Files:	/root/labs/Lab9.2
--------------------	-------------------

Step 1: Create the max_ordertotal View

1.1. In the previous lab, you defined a view named **max_ordertotal**. Use the **describe** command to verify:

```
hive> describe max_ordertotal;  
OK  
maxtotal      int      None  
userid        int      None
```

If you do not have this view, define it now as:

```
CREATE VIEW max_ordertotal AS  
SELECT max(ordertotal) AS maxtotal, userid  
FROM orders GROUP BY userid;
```

Step 2: Think in MapReduce

2.1. Consider the following join statement that you executed in the previous lab:

```
SELECT ordertotal, orders.userid, itemlist  
FROM orders  
JOIN max_ordertotal ON max_ordertotal.userid = orders.userid AND  
max_ordertotal.maxtotal = orders.ordertotal;
```


Recall this join statement required two MapReduce jobs to execute.

2.2. What if we could send all the orders by a particular customer to the same reducer? How could we accomplish this?

2.3. Suppose we have distributed the records so that we know the same reducer handles all orders from a customer. Then we could sort the orders by **totalorder** descending, and the first order would be their maximum order. Run the following query to understand the logic here:

```
SELECT * FROM orders DISTRIBUTE BY userid  
SORT BY userid, ordertotal DESC;
```

2.4. Look closely at the output. Each customer's largest order should appear first in his or her respective list of orders. For example, Caitlin F's largest order was \$600 on April 25, 2012:

72094	2012-04-25	100	Caitlin	F	600	...
87194	2013-01-05	100	Caitlin	F	588	...
53034	2011-06-11	100	Caitlin	F	588	...
56003	2011-07-30	100	Caitlin	F	588	...

The reducer gets all orders from a customer, and the first order the reducer receives is the largest one (which is what we are trying to find!). In the next step, you will use a custom reducer using Python that pulls off this top value.

Step 3: Use a Custom Reducer

3.1. Using a text editor, open the file **max_order.py** in the **Lab9.2** folder.

3.2. Notice this Python script prints the first line that it processes. Then it hangs on to the userid and skips all subsequent lines until the userid changes.

3.3. Copy max_order.py into the /tmp folder and make it executable:

```
# cp max_order.py /tmp  
# chmod +x /tmp/max_order.py
```

3.4. Start the Hive shell.

3.5. Add **max_order.py** as a resource using the **add file** command:

```
hive> add file /tmp/max_order.py; Added resource: max_order.py
```

NOTE: The **add file** command makes the file available to all mappers and reducers of this Hive query.

3.6. Specify three reducers so we can verify the logic of our query:

```
hive> set mapreduce.job.reduces=3;
```

3.7. Now run the following join query, which uses the Python script as its reducer. You may want to type this in a text file so you can rerun it easier if you have a typo, and make sure you use the proper path to **max_order.py**.

```
from (
select userid,ordertotal,itemlist from orders
distribute by userid
sort by userid,ordertotal DESC)
orders
insert overwrite directory 'maxorders' reduce userid,ordertotal,itemlist
using 'max_order.py';
```

The query should execute a single MapReduce job this time, and you should also notice three reducers.

Step 4: View the Results

4.1. From the command line, list the contents of the **maxorders** folder in HDFS. You should see three files, one from each reducer:

```
# hadoop fs -ls maxorders
Found 3 items
root hdfs      3611      maxorders/000000_0
root hdfs      3708      maxorders/000001_0
root hdfs      3714      maxorders/000002_0
```

4.2. View the contents of one of the files:

```
# hadoop fs -cat maxorders/000000_0
...

90588 Boots,Grill,Spark Plugs,Vacuum,Coffee Maker,DVD,2-Way
Radio,Dolls,Games,DVD,pillows,Pants
93600 Dishwasher,Table,Grill,DVD,DVD,DVD,Keychain,Dryer, Washer &
Dryer,Grill,Coffee Maker,pillows
96600 Table,Jeans,Washer,Wrench Set,Grill,Color Laser
Printer,Dryer,Air Compressor,DVD,Dolls,2-Way Radio,Sweater
99600 Washer,Cookware,Vacuum,Freezer,2-Way Radio,Bicycle,Washer &
Dryer,Coffee Maker,Refrigerator, DVD,Boots,DVD
```

The output shows the userid, ordertotal and itemlist of the largest order placed by each customer.

RESULT: You used a custom reducer (a Python script) to modify a Hive query that originally took two MapReduce jobs to execute so that it can now be executed in a single MapReduce job. You also learned how to assign a custom reducer (or mapper) to a Hive query.

ANSWERS:

2.2: Use the **DISTRIBUTE BY** clause and distribute the records by the **userid** column.